



Greedy Algorithms





A short list of categories

- Algorithm types we will consider include:
 - Simple recursive algorithms
 - Backtracking algorithms
 - Divide and conquer algorithms
 - Dynamic programming algorithms
 - ➔ ■ Greedy algorithms
 - Branch and bound algorithms
 - Brute force algorithms
 - Randomized algorithms



Optimization problems

- An **optimization problem** is one in which you want to find, not just *a* solution, but the *best* solution
- A “greedy algorithm” sometimes works well for optimization problems
- A **greedy algorithm** works in phases. At each phase:
 - You take the best you can get right now, without regard for future consequences
 - You hope that by choosing a *local* optimum at each step, you will end up at a *global* optimum



Example: Counting money

- Suppose you want to count out a certain amount of money, using the fewest possible bills and coins
- A greedy algorithm would do this would be:
At each step, take the largest possible bill or coin that does not overshoot
 - Example: To make \$6.39, you can choose:
 - a \$5 bill
 - a \$1 bill, to make \$6
 - a 25¢ coin, to make \$6.25
 - A 10¢ coin, to make \$6.35
 - four 1¢ coins, to make \$6.39
- For US money, the greedy algorithm always gives the optimum solution



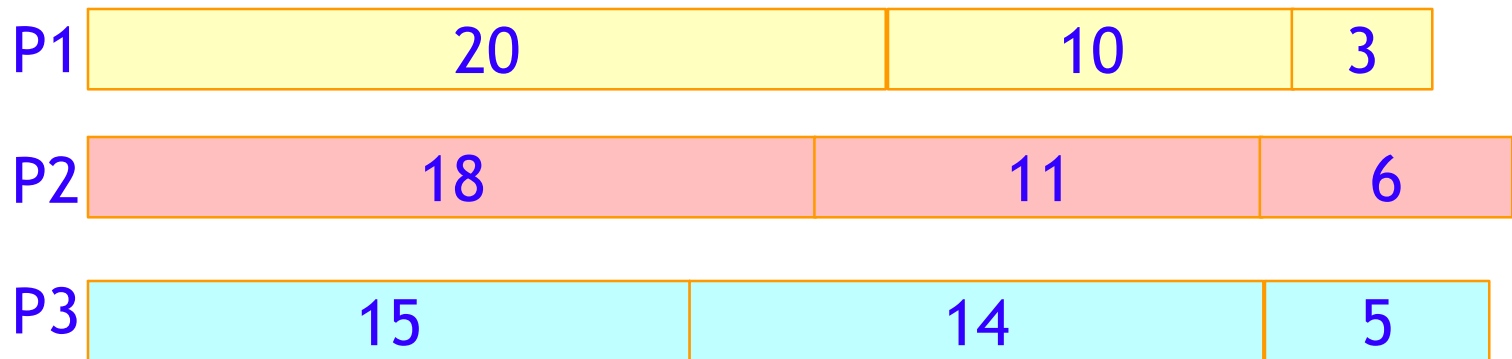
A failure of the greedy algorithm

- In some (fictional) monetary system, “krons” come in 1 kron, 7 kron, and 10 kron coins
- Using a greedy algorithm to count out 15 krons, you would get
 - A 10 kron piece
 - Five 1 kron pieces, for a total of 15 krons
 - This requires six coins
- A better solution would be to use two 7 kron pieces and one 1 kron piece
 - This only requires three coins
- The greedy algorithm results in a solution, but not in an optimal solution



A scheduling problem


- You have to run nine jobs, with running times of 3, 5, 6, 10, 11, 14, 15, 18, and 20 minutes
- You have three processors on which you can run these jobs
- You decide to do the longest-running jobs first, on whatever processor is available



- Time to completion: $18 + 11 + 6 = 35$ minutes
- This solution isn't bad, but we might be able to do better

Another approach

- What would be the result if you ran the *shortest* job first?
- Again, the running times are 3, 5, 6, 10, 11, 14, 15, 18, and 20 minutes

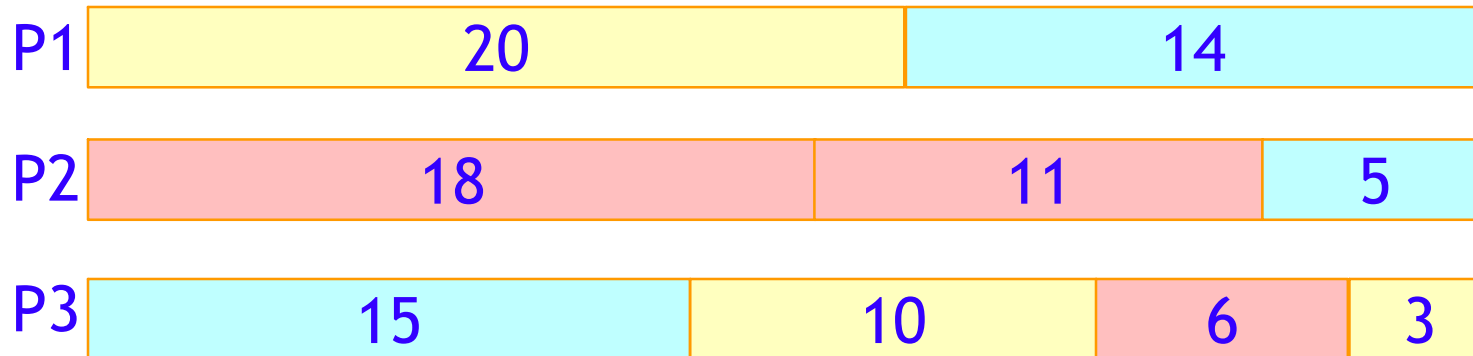


- That wasn't such a good idea; time to completion is now $6 + 14 + 20 = 40$ minutes
- Note, however, that the greedy algorithm itself is fast
 - All we had to do at each stage was pick the minimum or maximum



An optimum solution


- Better solutions do exist:



- This solution is clearly optimal (why?)
- Clearly, there are other optimal solutions (why?)
- How do we find such a solution?
 - One way: Try all possible assignments of jobs to processors
 - Unfortunately, this approach can take exponential time

Huffman encoding

- The Huffman encoding algorithm is a greedy algorithm
- You always pick the two smallest numbers to combine




A=00
B=100
C=01
D=1010
E=11
F=1011

- Average bits/char:
 $0.22*2 + 0.12*3 + 0.24*2 + 0.06*4 + 0.27*2 + 0.09*4 = 2.42$
- The Huffman algorithm finds an optimal solution

Minimum spanning tree


- A minimum spanning tree is a least-cost subset of the edges of a graph that connects all the nodes
 - Start by picking any node and adding it to the tree
 - Repeatedly: Pick any *least-cost* edge from a node in the tree to a node not in the tree, and add the edge and new node to the tree
 - Stop when all nodes have been added to the tree



- The result is a least-cost ($3+3+2+2+2=12$) spanning tree
- If you think some other edge should be in the spanning tree:
 - Try adding that edge
 - Note that the edge is part of a cycle
 - To break the cycle, you must remove the edge with the greatest cost
 - This will be the edge you just added

Traveling salesman

- A salesman must visit every city (starting from city **A**), and wants to cover the least possible distance
 - He can revisit a city (and reuse a road) if necessary
- He does this by using a greedy algorithm: He goes to the next nearest city from wherever he is



- From **A** he goes to **B**
- From **B** he goes to **D**
- This is *not* going to result in a shortest path!
- The best result he can get now will be **ABDBCE**, at a cost of **16**
- An actual least-cost path from **A** is **ADBCE**, at a cost of **14**



Analysis

- A greedy algorithm typically makes (approximately) n choices for a problem of size n
 - (The first or last choice may be forced)
- Hence the expected running time is:
 $O(n * O(\text{choice}(n)))$, where $\text{choice}(n)$ is making a choice among n objects
 - Counting: Must find largest useable coin from among k sizes of coin (k is a constant), an $O(k)=O(1)$ operation;
 - Therefore, coin counting is (n)
 - Huffman: Must sort n values before making n choices
 - Therefore, Huffman is $O(n \log n) + O(n) = O(n \log n)$
 - Minimum spanning tree: At each new node, must include new edges and keep them sorted, which is $O(n \log n)$ overall
 - Therefore, MST is $O(n \log n) + O(n) = O(n \log n)$



Other greedy algorithms

- Dijkstra's algorithm for finding the shortest path in a graph
 - Always takes the *shortest* edge connecting a known node to an unknown node
- Kruskal's algorithm for finding a minimum-cost spanning tree
 - Always tries the *lowest-cost* remaining edge
- Prim's algorithm for finding a minimum-cost spanning tree
 - Always takes the *lowest-cost* edge between nodes in the spanning tree and nodes not yet in the spanning tree



Dijkstra's shortest-path algorithm

- Dijkstra's algorithm finds the shortest paths from a given node to all other nodes in a graph
 - Initially,
 - Mark the given node as *known* (path length is zero)
 - For each out-edge, set the distance in each neighboring node equal to the *cost* (length) of the out-edge, and set its *predecessor* to the initially given node
 - Repeatedly (until all nodes are known),
 - Find an unknown node containing the smallest distance
 - Mark the new node as known
 - For each node adjacent to the new node, examine its neighbors to see whether their estimated distance can be reduced (distance to known node plus cost of out-edge)
 - If so, also reset the predecessor of the new node



Analysis of Dijkstra's algorithm I

- Assume that the *average* out-degree of a node is some constant k
 - Initially,
 - Mark the given node as *known* (path length is zero)
 - This takes $O(1)$ (constant) time
 - For each out-edge, set the distance in each neighboring node equal to the *cost* (length) of the out-edge, and set its *predecessor* to the initially given node
 - If each node refers to a list of k adjacent node/edge pairs, this takes $O(k) = O(1)$ time, that is, constant time
 - Notice that this operation takes *longer* if we have to extract a list of names from a hash table




Analysis of Dijkstra's algorithm II

- Repeatedly (until all nodes are known), (n times)
 - Find an unknown node containing the smallest distance
 - Probably the best way to do this is to put the unknown nodes into a priority queue; this takes $k * O(\log n)$ time *each* time a new node is marked “known” (and this happens n times)
 - Mark the new node as known -- $O(1)$ time
 - For each node adjacent to the new node, examine its neighbors to see whether their estimated distance can be reduced (distance to known node plus cost of out-edge)
 - If so, also reset the predecessor of the new node
 - There are k adjacent nodes (on average), operation requires constant time at each, therefore $O(k)$ (constant) time
 - Combining all the parts, we get:
 $O(1) + n*(k*O(\log n)+O(k))$, that is, $O(nk \log n)$ time

Connecting wires


- There are n white dots and n black dots, equally spaced, in a line
- You want to connect each white dot with some one black dot, with a minimum total length of “wire”
- Example:



- Total wire length above is $1 + 1 + 1 + 5 = 8$
- Do you see a greedy algorithm for doing this?
- Does the algorithm guarantee an optimal solution?
 - Can you prove it?
 - Can you find a counterexample?

Collecting coins

- A checkerboard has a certain number of coins on it
- A robot starts in the upper-left corner, and walks to the bottom left-hand corner
 - The robot can only move in two directions: right and down
 - The robot collects coins as it goes
- You want to collect *all* the coins using the *minimum* number of robots
- Example:



- Do you see a greedy algorithm for doing this?
- Does the algorithm guarantee an optimal solution?
 - Can you prove it?
 - Can you find a counterexample?



The End
